

CASE STUDY::Use Terraform to provision an AWS EC2 instance and an S3 bucket. Deploy a sample static website on the S3 bucket using the EC2 instance as the backend server.

My case study topic is as follows:-

Infrastructure as Code with Terraform

- Concepts Used: Terraform, AWS S3, and EC2.
- Problem Statement: "Use Terraform to provision an AWS EC2 instance and an S3 bucket. Deploy a sample static website on the S3 bucket using the EC2 instance as the backend server."
- Tasks:
 - Write a Terraform script to create an EC2 instance and an S3 bucket.
 - Deploy the static website on the S3 bucket.
 - Use the EC2 instance to interact with the S3 bucket and log the actions

Introduction to Infrastructure as Code (IaC) with Terraform

Infrastructure as Code (IaC) allows the automation of infrastructure deployment and management using code, offering efficiency, repeatability, and version control. Terraform is one of the most widely used IaC tools, allowing developers to define and manage infrastructure in cloud environments such as AWS, Google Cloud, and Azure. Terraform uses declarative configuration files that describe the desired state of the infrastructure, enabling consistent provisioning across environments. In this scenario, you are asked to use Terraform to provision two AWS services: an EC2 instance and an S3 bucket, and to deploy a static website on the S3 bucket while using the EC2 instance to interact with the bucket as a backend server.

Significance of Using Terraform in AWS

1. Automation: With Terraform, cloud resources can be created, modified, or destroyed automatically, reducing the need for manual intervention.
2. Version Control: The infrastructure code can be versioned and maintained in repositories, enabling teams to track changes, roll back configurations, and collaborate effectively.
3. Consistency: Infrastructure can be provisioned consistently across different environments, reducing configuration drift and ensuring that development, testing, and production environments are identical.
4. Multi-Cloud Support: Terraform provides the flexibility to manage resources not only on AWS but also across multiple cloud providers with a single tool.

Key Concepts and Services Involved are as follows:-

Terraform

- Significance: Terraform automates the provisioning and management of cloud resources. It allows organizations to treat their infrastructure as code, providing benefits like automation, collaboration, and consistency.

AWS EC2 (Elastic Compute Cloud)

- Significance: In this context, the EC2 instance will serve as the backend server that can interact with the S3 bucket. It will log the actions related to the deployment and management of the static website hosted on S3.
- Applications: EC2 is used for running web applications, backend services, databases, and other computing tasks. In this project, it plays a crucial role in managing and logging the interactions with the S3 bucket.

AWS S3 (Simple Storage Service)

- Significance: The S3 bucket will host the static website. S3 provides high availability and content delivery, making it a suitable platform for website hosting.
- Applications: AWS S3 is widely used for static website hosting, backup storage, big data analytics, and content distribution.

Static Website Hosting on S3

- Significance: Using S3 to host static websites is cost-effective and scalable. By leveraging this, the project eliminates the need for a dedicated web server, offloading the hosting duties to S3.
- Applications: This is useful for content-heavy websites, landing pages, and blogs that do not require server-side processing.

Applications of the Solution are as follows:-

1. Infrastructure Automation: The ability to automate infrastructure provisioning for websites, applications, or backend services significantly reduces deployment time and errors.
2. Scalable Website Hosting: Hosting static websites on S3 provides a scalable and highly available solution for website deployment. This is especially useful for static content, where high performance and low costs are critical.
3. Log Management and Monitoring: Using EC2 to interact with the S3 bucket for logging purposes provides deeper insight into website management. This setup can be extended to gather and analyze logs, track usage, or monitor performance metrics.
4. Collaboration and Versioning: By using Terraform, teams can collaborate effectively with shared infrastructure definitions, enabling version control of infrastructure similar to how they manage application code.

Step 1: Install AWS CLI

Specify user details

User details

User name

The user name can have up to 64 characters. Valid characters: A-Z, a-z, 0-9, and + = , . @ _ - (hyphen)

☒ Provide user access to the AWS Management Console - *optional*
If you're providing console access to a person, it's a [best practice](#) to manage their access in IAM Identity Center.

i Are you providing console access to a person?

User type

☐ Specify a user in Identity Center - Recommended
We recommend that you use Identity Center to provide console access to a person. With Identity Center, you can centrally manage user access to their AWS accounts and cloud applications.

☒ I want to create an IAM user
We recommend that you create IAM users only if you need to enable programmatic access through access keys, service-specific credentials for AWS CodeCommit or Amazon Keyspaces, or a backup credential for emergency account access.

Step 2: IAM user and access key creation

Console password

☐ Autogenerated password
You can view the password after you create the user.

☒ Custom password
Enter a custom password for the user.

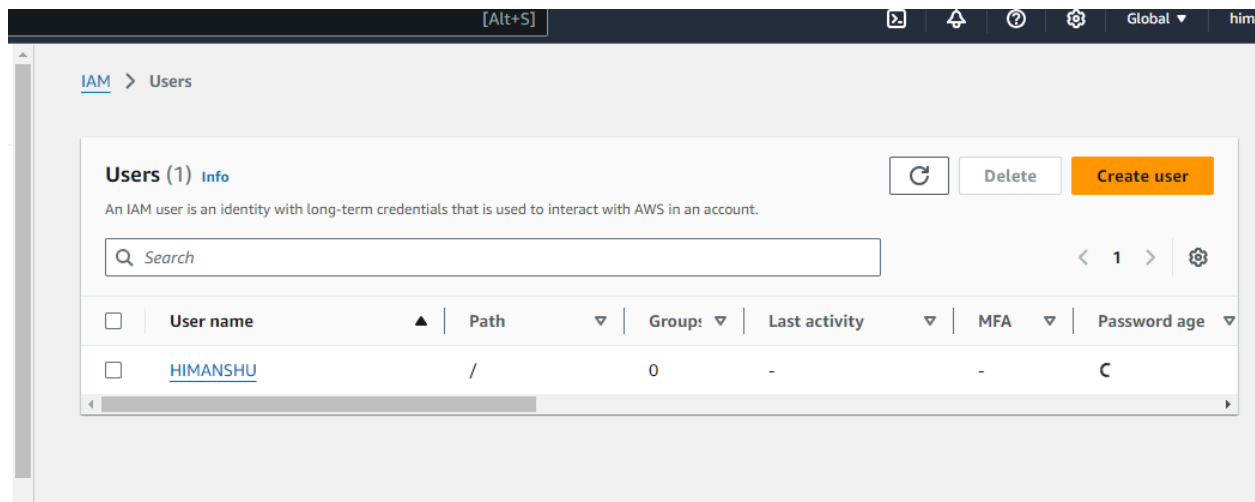
- Must be at least 8 characters long
- Must include at least three of the following mix of character types: uppercase letters (A-Z), lowercase letters (a-z), numbers (0-9), and symbols ! @ # \$ % ^ & * () _ + - (hyphen) = [] { } | ' "

☐ Show password

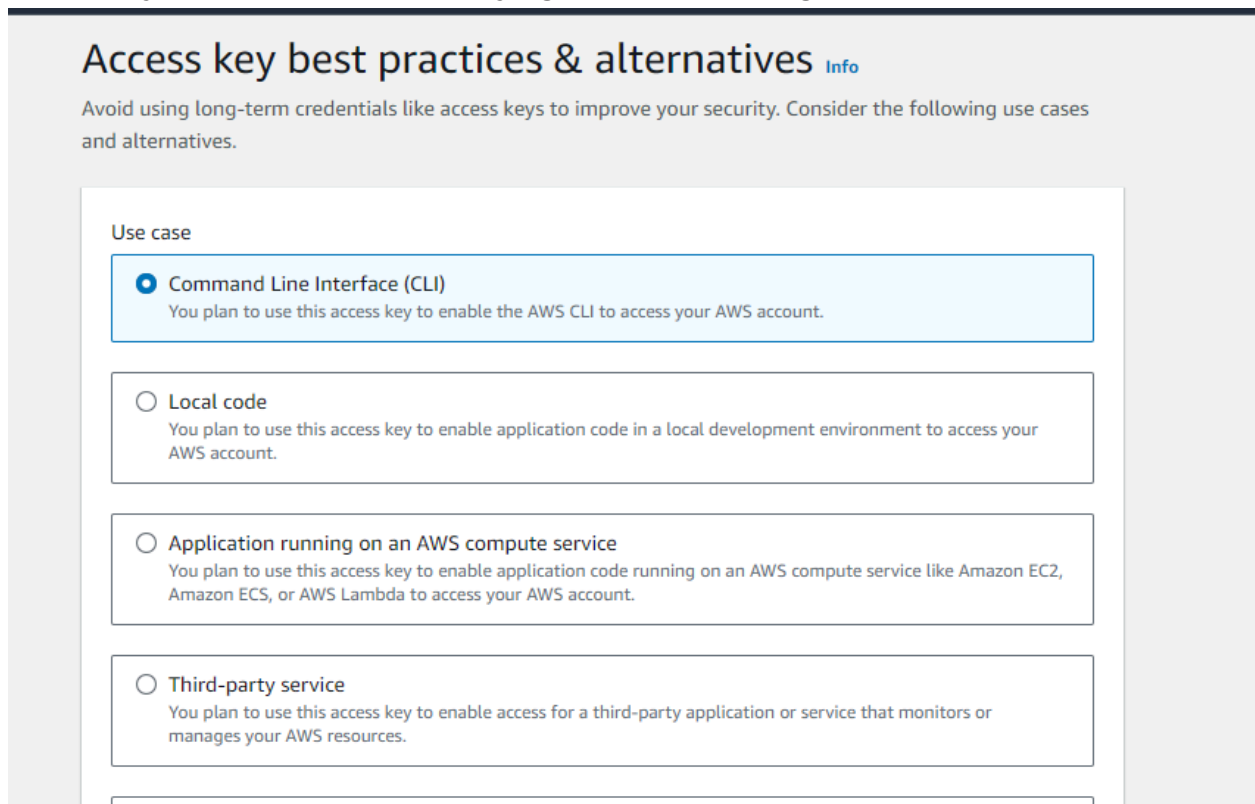
☒ Users must create a new password at next sign-in - Recommended
Users automatically get the [IAMUserChangePassword](#) policy to allow them to change their own password.

i If you are creating programmatic access through access keys or service-specific credentials for AWS CodeCommit or Amazon Keyspaces, you can generate them after you create this IAM user. [Learn more](#)

Cancel Next



Create a new Access key inside the IAM User by navigating to Security section. This access key would be asked while trying to run 'aws configure'.



Set permissions

Add user to an existing group or create a new one. Using groups is a best-practice way to manage user's permissions by job functions. [Learn more](#)

Permissions options

☐ Add user to group

Add user to an existing group, or create a new group. We recommend using groups to manage user permissions by job function.

☐ Copy permissions

Copy all group memberships, attached managed policies, and inline policies from an existing user.

☒ Attach policies directly

Attach a managed policy directly to a user. As a best practice, we recommend attaching policies to a group instead. Then, add the user to the appropriate group.

Permissions policies (1/1242)

Permissions policies (1/1242)

Choose one or more policies to attach to your new user.

Filter by Type

All types

< 1 2 3 4 5 6 7 ... 63 >

| <input type="checkbox"/> | Policy name | Type | Attached entities |
|-------------------------------------|--|----------------------------|-------------------|
| <input type="checkbox"/> | AccessAnalyzerServiceRole... | AWS managed | 0 |
| <input checked="" type="checkbox"/> | AdministratorAccess | AWS managed - job function | 0 |
| <input type="checkbox"/> | AdministratorAccess-Amplify | AWS managed | 0 |
| <input type="checkbox"/> | AdministratorAccess-AWSE... | AWS managed | 0 |
| <input type="checkbox"/> | AlexaForBusinessDeviceSet... | AWS managed | 0 |
| <input type="checkbox"/> | AlexaForBusinessFullAccess | AWS managed | 0 |
| <input type="checkbox"/> | AlexaForBusinessGatewayE... | AWS managed | 0 |
| <input type="checkbox"/> | AlexaForBusinessLifesized... | AWS managed | 0 |

Step 3: Terraform usage

Now, we are supposed to create a new Directory. I named it ADV DEVOPS CASE STUDY

Create 3 files in it ...

1. main.tf
2. Index.html
3. error.html

Now, initialize terraform in the directory by running terraform init command

```
Microsoft Windows [Version 10.0.22631.4317]
(c) Microsoft Corporation. All rights reserved.

C:\Users\2021h\OneDrive\Desktop\adv devops assignment>terraform init
Initializing the backend...
Initializing provider plugins...
- Finding latest version of hashicorp/aws...
- Finding latest version of hashicorp/tls...
- Finding latest version of hashicorp/local...
- Installing hashicorp/local v2.5.2...
- Installed hashicorp/local v2.5.2 (signed by HashiCorp)
- Installing hashicorp/aws v5.72.1...
- Installed hashicorp/aws v5.72.1 (signed by HashiCorp)
- Installing hashicorp/tls v4.0.6...
- Installed hashicorp/tls v4.0.6 (signed by HashiCorp)
Terraform has created a lock file .terraform.lock.hcl to record the provider
selections it made above. Include this file in your version control repository
so that Terraform can guarantee to make the same selections by default when
you run "terraform init" in the future.

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
```

Write this in index.html

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <meta name="viewport" content="width=device-width, initial-scale=1.0">
  <title>Welcome to My Static Website</title>
</head>
<body>
  <h1>Hello, World!</h1>
  <p>This is a sample static website hosted on S3.</p>
</body>
</html>
```

Write error.html like this

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <meta name="viewport" content="width=device-width, initial-scale=1.0">
```

```
<title>Error Page</title>
</head>
<body>
  <h1>Oops!</h1>
  <p>Sorry, the page you are looking for does not exist.</p>
</body>
</html>
```

Now, modify the code in main.tf like this.

This code contains

- 1. IAM user connection using access key id**
- 2. EC2 instance and bucket configuration**
- 3. Code snippets for connecting index.html and error.html**
- 4. Code snippets for creating a separate public and private key**

AWS Provider

```
provider "aws" {
  region = "us-east-1" # Use the region you configured with `aws configure`
}
```

IAM Role for EC2 to Access S3

```
resource "aws_iam_role" "ec2_role" {
  name = "ec2_role"
```

```
  assume_role_policy = jsonencode({
    Version: "2012-10-17",
    Statement: [{
      Action: "sts:AssumeRole",
      Effect: "Allow",
      Principal: {
        Service: "ec2.amazonaws.com"
      }
    }]
  })
}
```

Attach S3 Full Access Policy to the EC2 Role

```
resource "aws_iam_policy_attachment" "ec2_s3_access" {
  name      = "ec2-s3-access"
  roles     = [aws_iam_role.ec2_role.name]
  policy_arn = "arn:aws:iam::aws:policy/AmazonS3FullAccess"
}
```

Instance Profile for EC2

```
resource "aws_iam_instance_profile" "ec2_profile" {
  name = "ec2_profile"
  role = aws_iam_role.ec2_role.name
}

# Generate an SSH Key Pair
resource "tls_private_key" "my_key" {
  algorithm = "RSA"
  rsa_bits  = 2048
}

resource "aws_key_pair" "my_key" {
  key_name   = "id_rsa_himanshu_05122004"
  public_key = tls_private_key.my_key.public_key_openssh
}

# Save the private key to a file (for local access)
resource "local_file" "private_key" {
  filename = "${path.module}/id_qwe" # Adjust the path if necessary
  content  = tls_private_key.my_key.private_key_pem
}

# EC2 Instance
resource "aws_instance" "example" {
  ami          = "ami-06b21ccaef8cd686" # Change this AMI based on your region if
  necessary
  instance_type = "t2.micro"
  iam_instance_profile = aws_iam_instance_profile.ec2_profile.name
  key_name        = aws_key_pair.my_key.key_name # Associate the key pair with the
  instance

  tags = {
    Name = "advance devops"
  }
}

# S3 Bucket
resource "aws_s3_bucket" "website_bucket" {
  bucket = "himanshu18102003advdevops"
}

resource "aws_s3_bucket_public_access_block" "website_bucket_public_access" {
  bucket = aws_s3_bucket.website_bucket.id
}
```



```
    block_public_acls    = false
    ignore_public_acls   = false
    block_public_policy  = false
    restrict_public_buckets = false
  }

  resource "aws_s3_bucket_website_configuration" "website_config" {
    bucket = aws_s3_bucket.website_bucket.id

    index_document {
      suffix = "index.html"
    }

    error_document {
      key = "error.html"
    }
  }

  # Upload files to the S3 bucket
  resource "aws_s3_object" "website_index" {
    bucket = aws_s3_bucket.website_bucket.bucket
    key    = "index.html"
    source = "index.html"
  }

  resource "aws_s3_object" "website_error" {
    bucket = aws_s3_bucket.website_bucket.bucket
    key    = "error.html"
    source = "error.html"
  }
```

So, after commenting out the **acl attribute** in bucket and both html files, terraform plan and apply commands started working correctly

Terraform plan output

```

+ create

Terraform will perform the following actions:

# aws_instance.example will be created
+ resource "aws_instance" "example" {
  + ami                     = "ami-00b21ccaefffcdd00e"
  + arm                    = (known after apply)
  + associate_public_ip_address = (known after apply)
  + availability_zone       = (known after apply)
  + cpu_core_count          = (known after apply)
  + cpu_threads_per_core    = (known after apply)
  + disable_api_stop        = (known after apply)
  + disable_api_termination = (known after apply)
  + ebs_optimized           = (known after apply)
  + get_password_data       = false
  + host_id                 = (known after apply)
  + host_resource_group_arn = (known after apply)
  + iam_instance_profile    = "ec2_profile"
  + id                      = (known after apply)
  + instance_initiated_shutdown_behavior = (known after apply)
  + instance_lifecycle      = (known after apply)
  + instance_state          = (known after apply)
  + instance_type           = "t2.micro"
  + ipv6_address_count      = (known after apply)
  + ipv6_addresses          = (known after apply)
  + key_name                = "id_rsa_rohit_05122004"
  + ... (known after apply)
}

```

Terraform apply output

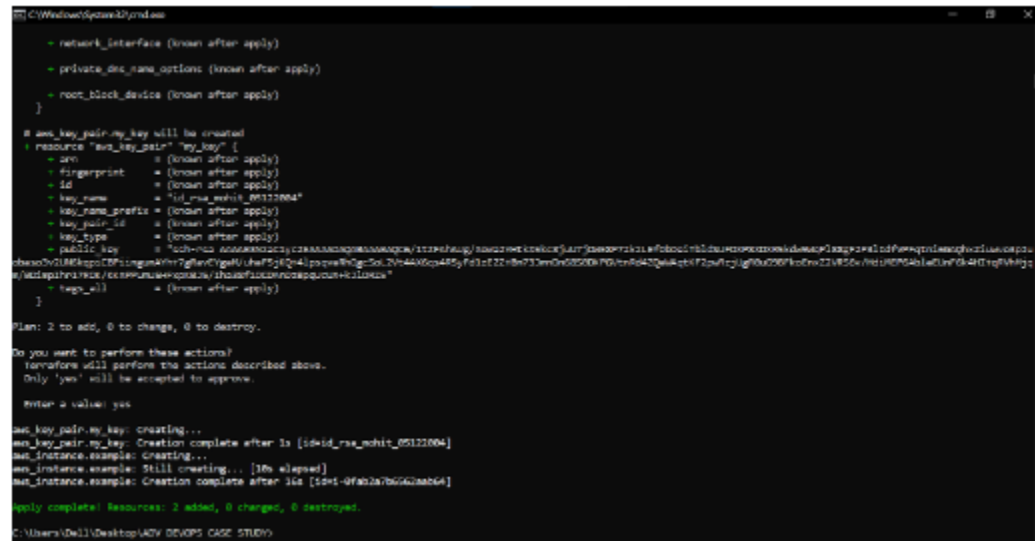
```

[C:\Windows\System32\cmd.exe]
+ aws_block_device (known after apply)
+ enclave_options (known after apply)
+ ephemeral_block_device (known after apply)
+ instance_market_options (known after apply)
+ maintenance_options (known after apply)
+ metadata_options (known after apply)
+ network_interface (known after apply)
+ private_dns_name_options (known after apply)
+ root_block_device (known after apply)
}

# aws_key_pair.my_key will be created
+ resource "aws_key_pair" "my_key" {
  + arm                    = (known after apply)
  + fingerprint            = (known after apply)
  + id                    = (known after apply)
  + key_name               = "id_rsa_rohit_05122004"
  + key_name_prefix        = (known after apply)
  + key_pair_id            = (known after apply)
  + key_type               = (known after apply)
  + public_key             = "ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQBAQC1L1P3H4uJ/RwW2H8U5C8JwA7/DwEP7Zr1f5F0xwL7oLdRFD7B2C83FubwWw*188y*2F012rP7Fqtn[68]Qn2i1wv8p3
  + ... (known after apply)
  + tags_all               = (known after apply)
}

plan: 2 to add, 0 to change, 0 to destroy.

```



Objects (2) Info

Actions ▾

Create folder

Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

< 1 >

| <input type="checkbox"/> | Name | Type | Last modified | Size | Storage class |
|--------------------------|----------------------------|------|--|---------|---------------|
| <input type="checkbox"/> | error.html | html | October 18, 2024, 01:26:11 (UTC+05:30) | 297.0 B | Standard |
| <input type="checkbox"/> | index.html | html | October 18, 2024, 01:26:11 (UTC+05:30) | 517.0 B | Standard |

The ec2 instance has also been created ..
Now, we will use it to interact with our S3 bucket

User details

| | | |
|-----------------------|--|-------------------------------|
| User name HIMANSHU | Console password type Custom password | Require password reset Yes |
|-----------------------|--|-------------------------------|

Permissions summary

< 1 >

| Name | Type | Used as |
|---------------------------------------|----------------------------|--------------------|
| AdministratorAccess | AWS managed - job function | Permissions policy |
| IAMUserChangePassword | AWS managed | Permissions policy |

Tags - optional

Tags are key-value pairs you can add to AWS resources to help identify, organize, or search for resources. Choose any tags you want to associate with this user.

No tags associated with the resource.

Add new tag

Services

Search

Global

Himanshu

User created successfully

You can view and download the user's password and email instructions for signing in to the AWS Management Console.

Step 1
Specify user details

Step 2
Set permissions

Step 3
Review and create

Step 4
Retrieve password

Retrieve password

You can view and download the user's password below or email users instructions for signing in to the AWS Management Console. This is the only time you can view and download this password.

Console sign-in details

Email sign-in instructions

Console sign-in URL
 <https://440744244657.signin.aws.amazon.com/console>

User name
 HIMANSHU

Console password
 ***** [Show](#)

Cancel

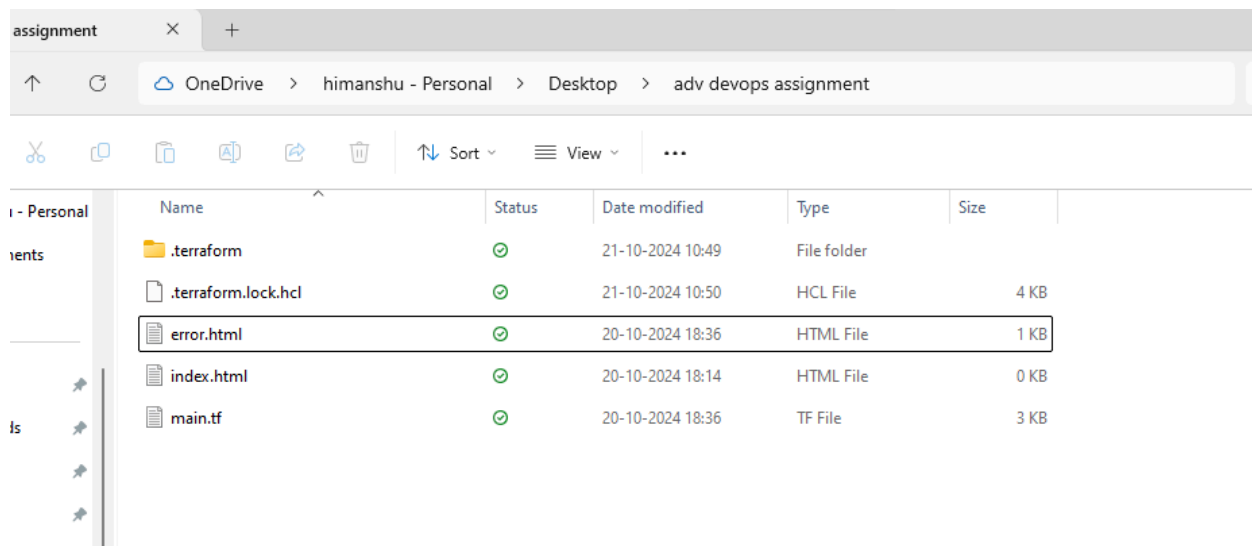
Download .csv file

Return to users list

In our main.tf code, we have mentioned two code blocks where a public and a private key is to be created. So that accessing our EC2 instance should not be problematic
Public key "id_qwe_himanshu_05122004" is shown to have been created in AWS key pair section

| <input type="checkbox"/> | Name | Type | Created | Fingerprint | ID |
|--------------------------|--------------------------|------|---------------------------|---------------------------------|-----------------------|
| <input type="checkbox"/> | id_qwe_himanshu_05122005 | rsa | 2024/10/21 09:40 GMT+5:30 | e7:53:16:5d:a5:c2:e4:92:f9:d... | key-035b82bceca4d395b |

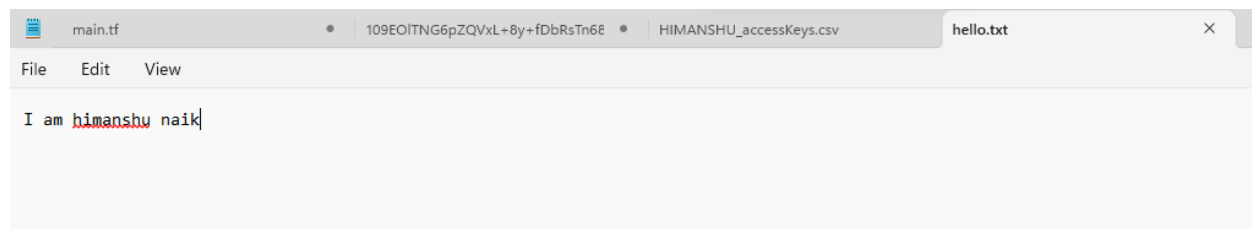
private key "id_rsa" is created in the same directory as our main.tf file



Use this command to remotely login onto your ec2 instance
ssh -i "<path to your main.tf file/ has to be the same path>\id_rsa"

ec2-user@ec2-54-236-238-111.compute-1.amazonaws.com

MAKE A TEXT FILE

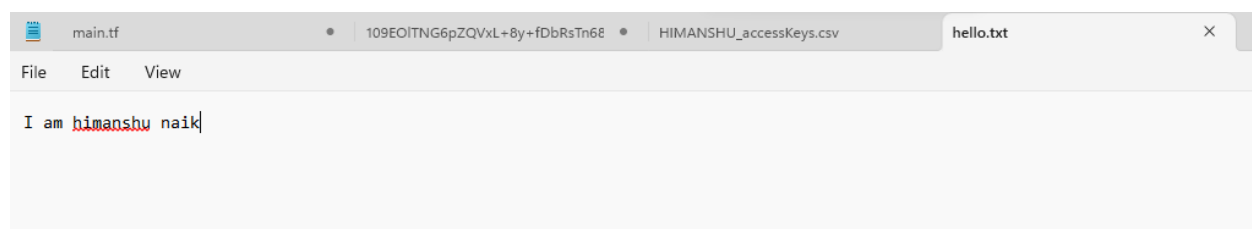
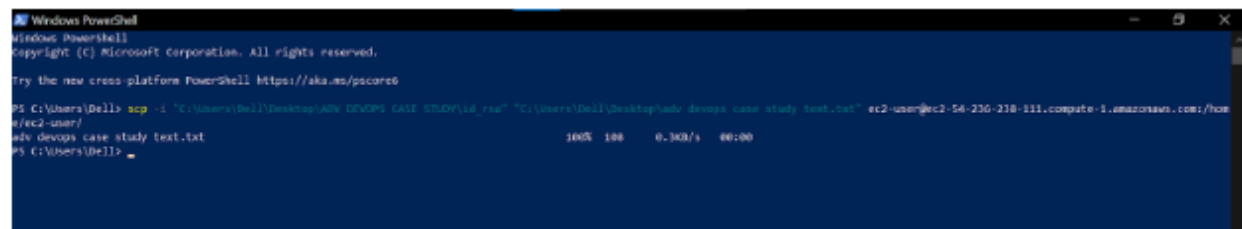


Open new terminal and run this command to upload our newly made text file on our instance

scp -i "path_to_your_key.pem" "path_to_local_file"

ec2-user@your-ec2-public-ip:/home/ec2-user/.

This command is used to upload a random text file onto our instance



Objects (3) [Info](#)

[Refresh](#)
[Copy S3 URI](#)
[Copy URL](#)
[Download](#)
[Open](#)
[Delete](#)
[Actions](#)

[Create folder](#)
[Upload](#)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

| <input type="checkbox"/> | Name | Type | Last modified | Size | Storage class |
|--------------------------|----------------------------|------|--|---------|---------------|
| <input type="checkbox"/> | error.html | html | October 21, 2024, 09:40:55 (UTC+05:30) | 297.0 B | Standard |
| <input type="checkbox"/> | hello.txt | txt | October 21, 2024, 10:07:06 (UTC+05:30) | 18.0 B | Standard |
| <input type="checkbox"/> | index.html | html | October 21, 2024, 09:40:55 (UTC+05:30) | 317.0 B | Standard |

Then upload that file onto our bucket from the original terminal and list them just for confirmation

Command to upload the file onto our bucket: **aws s3 cp /home/ec2-user/you.txt**

s3://your-s3-bucket-name/

```
Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

Install the latest PowerShell for new features and improvements! https://aka.ms/PSWindows

PS C:\Users\2021h> scp
usage: scp [-346ABCOpgRrsTv] [-c cipher] [-D sftp_server_path] [-F ssh_config]
        [-i identity_file] [-J destination] [-l limit] [-o ssh_option]
        [-P port] [-S program] [-X sftp_option] source ... target
PS C:\Users\2021h> scp -i "C:/Users/2021h/OneDrive/Desktop/adv devops case study/id_qwe" "C:/Users/2021h/OneDrive/Desktop/hello.txt" ec2-user@ec2-3-89-67-93.compute-1.amazonaws.com:~/
hello.txt
100% 18 0.0KB/s 00:00
PS C:\Users\2021h>
PS C:\Users\2021h>
```

Use this command in the original terminal to log the actions of the S3 bucket.

echo "Uploaded file 'filename.txt' to S3 on \$(date)" >> action_log.txt

```
ctory.
ec2-user@ec2-3-89-67-93.compute-1.amazonaws.com: Permission denied (publickey,gssapi-keyex,gssapi-with-mic).
Connection closed
[ec2-user@ip-172-31-39-32 ~]$ aws s3 cp /home/ec2-user/hello.txt s3://himanshu18102003advdevops/
upload: ./hello.txt to s3://himanshu18102003advdevops/hello.txt
[ec2-user@ip-172-31-39-32 ~]$ ls
hello.txt
[ec2-user@ip-172-31-39-32 ~]$ aws s3 ls
2024-10-10 21:06:05 exp12d15c63
2024-10-21 04:10:48 himanshu18102003advdevops
2024-10-10 20:09:41 himanshunaik123
2024-10-10 20:39:27 naikhim
[ec2-user@ip-172-31-39-32 ~]$ echo "Uploaded file 'hello.txt' to S3 on $(date)" >> action_log.txt
[ec2-user@ip-172-31-39-32 ~]$ cat action_log.txt
Uploaded file 'hello.txt' to S3 on Mon Oct 21 04:38:26 UTC 2024
[ec2-user@ip-172-31-39-32 ~]$ echo "Uploaded file 'hello.txt' to S3 on $(date)" >> action_log.txt_
```

Conclusion: Thus, by creating an IAM user, providing an access key to it , installing AWS CLI, using terraform commands in our project to create ec2 instance and s3 bucket, we were successfully able to ensure interaction between our ec2 instance and s3 bucket. We tested this by uploading a text file over to our s3 bucket by using our ec2 instance. This experiment ensured that terraform can be used to change the configuration of our resources and create and access the proper ones using the AWS CLI